

Day4 - Markov chain Monte Carlo methods

Michel Bje Randahl Nielsen (s093481)

July 1, 2015

1 About the implementations...

All implementations are done in F# scripts but uses functionality offered by the R programming environment. The scripts utilizes the type provider functionality which F# provides, to invoke functions in R.

F# project page: <http://fsharp.org/>

F# RProvider project page: <https://bluemountaincapital.github.io/FSharpRProvider/>

All scripts has been developed inside visual studio with F# 3.1, but is only dependend on the existance of .Net 4.5 and a F# 3.1 instalation to be run. A full project has been attached to these hand ins and contains all of the scripts and nescesary dll's.

Note that all scripts must run following code in the beginning in order to function, but this piece of code has been omitted in these reports.


```
#r @"..\packages\R.NET.Community.1.5.16\lib\net40\RDotNet.dll"
#r @"..\packages\R.NET.Community.FSharp.0.1.9\lib\net40\RDotNet.FSharp.dll"
#r @"..\packages\RProvider.1.1.8\lib\net40\RProvider.dll"
#r @"..\packages\RProvider.1.1.8\lib\net40\RProvider.Runtime.dll"

open RDotNet
open RProvider
open RProvider.stats
open RProvider.graphics
open RProvider.mvtnorm
open RProvider.MASS
open System
```

2 Exercises

- 2.1 Simulate a Markov chain with values in $S = \{0,1\}$ such that $P(x_i = 1 - x_{i-1} = 0) = 1/\sqrt{2}$ and $P(x_i = 1 - x_{i-1} = 1) = 1/$ Plot the result.

```
let markov_chain n =  
  //defining the transition function for the markov chain  
  let transition_fun = function  
    //return 1 with a probability of 1/sqrt(2) if input is 0  
    | 0 -> R.rbinom(1, 1, 1.0 / Math.Sqrt(2.0)).AsNumeric()  
    |> Seq.head  
    |> int  
    //return 1 with a probability of 1/PI if input is 1  
    | _ -> R.rbinom(1, 1, 1.0 / Math.PI).AsNumeric()  
    |> Seq.head  
    |> int  
  
  //recursive loop to perform n-number of iterations of the markov chain  
  let rec loop n' prev = seq {  
    match n' with  
    //end case  
    | 0 -> yield transition_fun prev  
    //loop case  
    | _ -> let res = transition_fun prev  
            yield res  
            yield! loop (n'-1) res  
  }  
  
  //initializing the loop  
  R.rbinom(1, 1, 0.5).AsNumeric()  
  |> Seq.head  
  |> int  
  |> loop n  
  
  //running the markov chain and plotting the result in a histogram  
  markov_chain 1000  
  |> R.hist
```



The probability of ending up in a state of 0 is 0.51 and the probability of ending up in a state of 1 is 0.49.

2.2 Simulate a Markov chain with the Metropolis algorithm with the Cauchy distribution $C(0,1)$ as a target distribution. Use e.g. a normal pdf with 0 mean and a unit variance as a proposal q .

Defining the markov chain function, transition function, auxillary function based on normal distribution, target function using the cauchy distribution and the iteration loop. The metropolis hastings algorithm works by rejecting or accepting a sequence of random values from an auxillary distribution (in this case normal dist). The rejection is done as a basis of the previous accepted value which is used in the calculaton of the metropolis ratio that determines the probability for which a proposed value should be accepted or not.

```
let markov_chain n =
  //target function
  let f (x: float) =
    R.dcauchy(x, 0, 1).AsNumeric()
    |> Seq.head

  //auxillary function
  let q (x: float) =
    R.dnorm(x, mean=0, sd=1).AsNumeric()
    |> Seq.head

  //transition function using the basic metropolis algorithm
  let transition_fun prev =
    //candidate value
    let y = R.rnorm(1, mean=0, sd=1).AsNumeric()
    |> Seq.head

    //compute metropolis ratio
    let met_ratio = f y * q prev / (f prev * q y)

    let accept_prob = Seq.min [1.0; met_ratio]
```

```

//toss unfair coin with accept prob
let accept = R.rbinom(1, 1, accept_prob).AsNumeric()
      |> Seq.head
      |> int
match accept with
| 0 -> prev
| _ -> y

//recursive loop to perform n' iterations
let rec loop n' prev = seq {
  match n' with
    //end case
  | 0 -> yield transition_fun prev
    //loop case
  | _ -> let res = transition_fun prev
        yield res
        yield! loop (n' - 1) res
}

//initializing the loop
R.rnorm(1, mean=0, sd=1).AsNumeric()
|> Seq.head
|> loop n

```

```

let n = 10000
let result = markov_chain n |> List.ofSeq

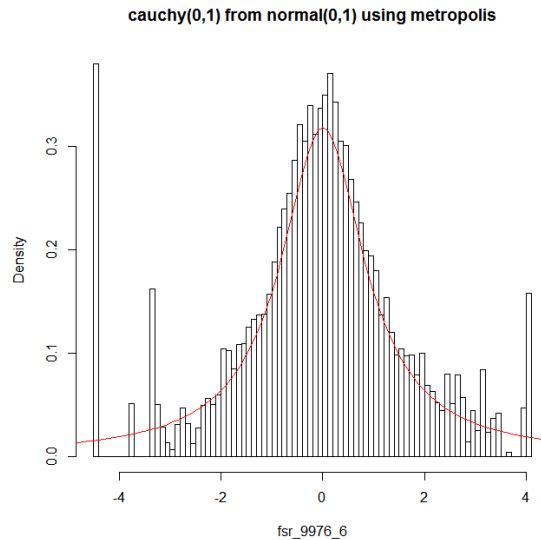
```

Plotting the data in a histogram.

```

namedParams [
  "x", box result
  "breaks", box 100
  "main", box "cauchy(0,1) from normal(0,1) using metropolis"
  "probability", box true
]
|> R.hist
let xs = [-5.0 .. 0.1 .. 5.0]
let normal_fun =
  xs
  |> List.map (fun x -> Seq.head(R.dcauchy(x, 0, 1).AsNumeric()))
namedParams [
  "x", box xs
  "y", box normal_fun
  "col", box "red"
]
|> R.lines

```



It can be seen that the resulting distribution follows the cauchy dist in the center, but fails to follow the dist in the extremes.

2.3 Same as previous exercise with a Metropolis-Hastings algorithm with Gaussian proposal. How does the simulated Markov chain depend on the variance of the proposal distribution?

```
//the difference between metropolis-hastings and the basic metropolis
//algorithms are that the metropolis-hastings algorithm calculates the
//auxillary function based on the proposed state and the previous state,
//and that the proposed value is drawn from a distribution with a mean
//based on the previous accepted value
let markov_chain n =
  //target function
  let f (x: float) =
    R.dcauchy(x, 0, 1).AsNumeric()
    |> Seq.head

  //auxillary function
  let q (x: float) (y: float) =
    R.dnorm(x, mean=y, sd=1).AsNumeric()
    |> Seq.head

  //transition function using the metropolis-hastings algorithm
  let transition_fun prev =
    //candidate value
    let y = R.rnorm(1, mean=prev, sd=1).AsNumeric()
    |> Seq.head

    //compute metropolis ratio
    let met_ratio = f y * q prev y / (f prev * q y prev)

    let accept_prob = Seq.min [1.0; met_ratio]
    //toss unfair coin with accept prob
    let accept = R.rbinom(1, 1, accept_prob).AsNumeric()
```

```

        |> Seq.head
        |> int
    match accept with
    | 0 -> prev
    | _ -> y

//recursive loop to perform n iterations in the markov chain
let rec loop n' prev = seq {
    match n' with
        //end case
    | 0 -> yield transition_fun prev
        //loop case
    | _ -> let res = transition_fun prev
            yield res
            yield! loop (n' - 1) res
}

R.rnorm(1, mean=0, sd=1).AsNumeric()
|> Seq.head
|> loop n

```

```

let n = 25000
let result = markov_chain n |> List.ofSeq

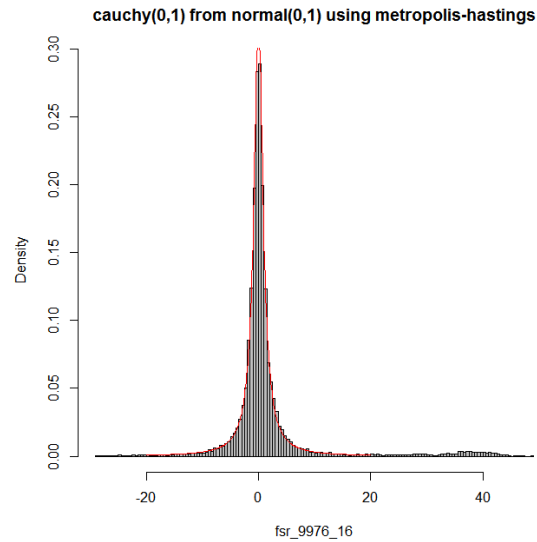
```

Plotting the data in a histogram.

```

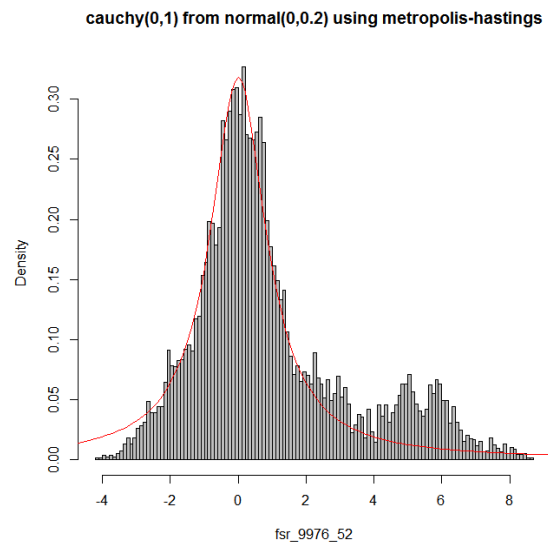
namedParams [
    "x", box result
    "breaks", box 150
    "main", box "cauchy(0,1) from normal(0,1) using metropolis-hastings"
    "col", box "grey"
    "probability", box true
]
|> R.hist
let xs = [-20.0 .. 0.1 .. 20.0]
let normal_fun =
    xs
    |> List.map (fun x -> Seq.head(R.dcauchy(x, 0, 1).AsNumeric()))
namedParams [
    "x", box xs
    "y", box normal_fun
    "col", box "red"
]
|> R.lines

```



The histogram reveals that the sampled data has reached its target distribution very well.

To get a sense of the effect that the variance has on the result, we generate a sample using an auxiliary function based on the normal dist but with a very low variance.



Here it can be seen that a very low variance will result in a skewed distribution that doesn't follow the target distribution.

2.4 Simulate a sample from a Beta $B(1/2, 1/2)$ with the rejection algorithm (using e.g. the uniform distribution as majorating density). Then simulate the same $B(1/2, 1/2)$ with the MH algorithm.

First simulating using the rejection method.

```
//simulating using the rejection algorithm
let n = 25000
let U1 = R.runif(n).AsNumeric()
let U2 = R.runif(n).AsNumeric()

let accept_function scaling_factor (u1, u2) =
```



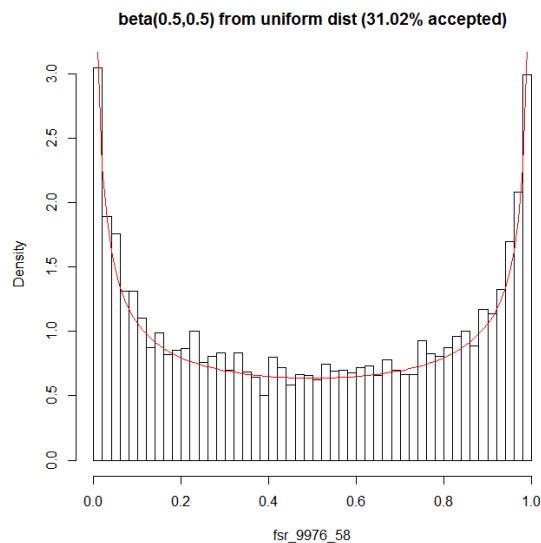
```

let beta_res =
  R.dbeta(u1, shape1=0.5, shape2=0.5).AsNumeric()
  |> Seq.head
scaling_factor * u2 < beta_res

let rbeta_rej =
  Seq.zip U1 U2
  |> Seq.filter (accept_function 3.0)
  |> Seq.map fst
let accepted = 100.0 * float(rbeta_rej |> Seq.length) / float n

//plotting the result of the rejection algorithm
namedParams [
  "x", box rbeta_rej
  "main", box (sprintf "beta(0.5,0.5) from uniform dist (%A%% accepted)"
    accepted)
  "breaks", box 50
  "probability", box true
]
|> R.hist
let xs = [0.0 .. 0.01 .. 1.0]
namedParams [
  "x", box xs
  "y", box <| R.dbeta(xs, 0.5, 0.5)
  "col", box "red"
]
|> R.lines

```



Next simulating using the MH algorithm.

```

//simulating using the MH algorithm
let markov_chain n =
  //target function
  let f (x: float) =
    R.dbeta(x, 0.5, 0.5).AsNumeric()

```

```

|> Seq.head

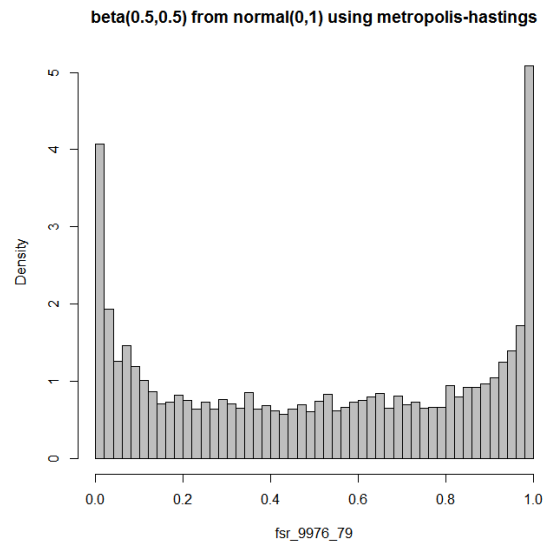
//auxillary function
let q (x: float) (y: float) =
  R.dnorm(x, mean=y, sd=1).AsNumeric()
|> Seq.head

//transition function based on MH algorithm
let transition_fun prev =
  let y = R.rnorm(1, mean=prev, sd=1).AsNumeric()
    |> Seq.head
  let met_ratio =
    f y * q prev y /
    (f prev * q y prev)
  let accept_prob = Seq.min [1.0; met_ratio]
  let accept = R.rbinom(1, 1, accept_prob).AsNumeric()
    |> Seq.head
    |> int
  match accept with
  | 0 -> prev
  | _ -> y
let rec loop n' prev = seq {
  match n' with
  | 0 -> yield transition_fun prev
  | _ ->
    let res = transition_fun prev
    yield res
    yield! loop (n' - 1) res
}

R.rnorm(1, mean=0, sd=1).AsNumeric()
|> Seq.head
|> loop n

//calculating and plotting the result
let rbeta_mh = markov_chain n |> List.ofSeq
namedParams [
  "x", box <| Seq.filter (fun x -> x >= 0.0) rbeta_mh
  "breaks", box 50
  "main", box "beta(0.5,0.5) from normal(0,1) using metropolis-hastings"
  "col", box "grey"
  "probability", box true
]
|> R.hist
namedParams [
  "x", box xs
  "y", box <| R.dbeta(xs, 0.5, 0.5)
  "col", box "red"
]
|> R.lines

```




based on the histograms alone, it is difficult to spot any difference between the two results.

- Plotting the autocorrelation.

```
//plotting the autocorrelation
namedParams [
  "mfrow", [1;2]
]
|> R.par

namedParams [
  "x", box rbeta_rej
  "main", box "rejection method"
]
|> R.acf

namedParams [
  "x", box rbeta_mh
  "main", box "MH algo"
]
|> R.acf
```



from the autocorrelation plots, it is obvious that the MH algorithm produces data that is highly autocorrelated

and the rejection algorithm does not produce autocorrelated data. The autocorrelation seen in the MH generated data, is clearly due to the nature of markov processes where the state is dependend on the previous state.

2.5 Simulate a sample of size 1000 using the Gibbs sampler

simulating a MVN sample with mean 0 and covariance matrix with $\rho=0.7$ using the gibbs sampler.

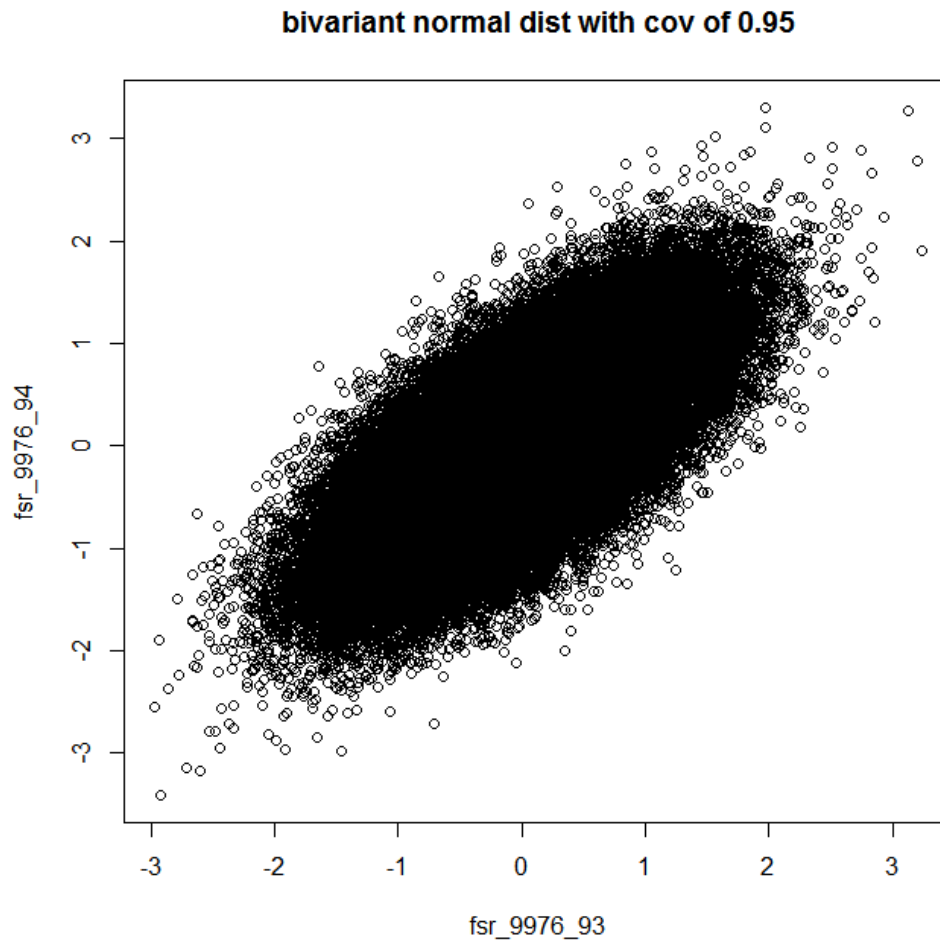
```
let n = 100000
//choosing a rho of 0.7
let rho = 0.7

let gibbs_sampler n rho =
  //gibbs sampling function
  let gibbs_sample x' =
    let y = R.rnorm(1, rho*x', 1.0 - rho**2.0 ).AsNumeric()
    |> Seq.head
    //generate x using y to adjust the mean
    let x = R.rnorm(1, rho*y, 1.0 - rho**2.0 ).AsNumeric()
    |> Seq.head
    x,y

  //recursive loop for the gibbs sampling process
  let rec loop n' x' = seq {
    match n' with
    //end case
    | 0 -> yield gibbs_sample x'
    //loop case
    | _ -> let res = gibbs_sample x'
    yield res
    yield! loop (n'-1) (fst res)
  }

  //initialize loop
  R.rnorm(1, 0, 1).AsNumeric()
  |> Seq.head
  |> loop n

let result = gibbs_sampler n rho
  |> List.ofSeq
  |> List.unzip
result
||> fun X Y ->
  namedParams [
    "x", box X
    "y", box Y
    "main", box "bivariant normal dist with cov of 0.95"
  ]
|> R.plot
```



Testing that the covariance matrix of the sample data is similar to the target covariance.

```
//testing the covariance matrix
//if the relationships in the covariance matrix matches the relationships
  from the original covariance matrix, then we have succeeded the
  simulation
let cov =
  result
  ||> fun X Y ->
    array2D [X ; Y]
  |> R.t
  |> fun x -> R.cov(x).AsNumericMatrix().ToArray()

//calculate the rho value
let rho' = (cov.[0,1] + cov.[1,0]) / (cov.[0,0] + cov.[1,1])
```

Outputs...

```
val cov : float [,] = [[0.5070249569; 0.3543349309]
                        [0.3543349309; 0.5078870188]]
val rho' : float = 0.6982574636
```

The estimated rho' value is similar to the target rho value and the proportions of the covariance matrix is similar to the target covariance matrix.